

# Linear Regression

Hongchang Gao

Spring 2024

# Summarization

- 1. Build model

$$f(x_i) = w_0 + w_1 x_{i,1} + w_2 x_{i,2} + \cdots + w_d x_{i,d}$$

- 2. Build loss function

$$\min_w \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$$

- 3. Optimize model

$$w = (X^T X)^{-1} X^T y$$

- 4. Make prediction

$$\hat{y} = w^T x_{n+1}$$

# Gradient Descent


- Machine learning model is an optimization problem

$$\min_x f(x)$$

- $f(x)$  is the objective function, or loss function.
- $x$  is the model parameter.
- Goal: find a model parameter to minimize the objective function

# Gradient Descent

- Gradient Descent method:
  - Decrease the loss function along the direction of the negative gradient
    - The gradient represents the slope of the loss function curve



# Gradient Descent

- Gradient Descent method:
  - Decrease the loss function along the direction of the negative gradient

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

- $x_t$  is the model parameter in the t-th iteration
- $\nabla f(x_t)$  is the gradient
- $\eta$  is the learning rate, or step size

# Gradient Descent

- Example: Linear Regression

$$\min_w \|y - Xw\|^2$$

- Gradient:

$$-2X^T(y - Xw)$$

- Gradient descent:

$$w_{t+1} = w_t - \eta(2X^T(Xw_t - y))$$

# Example

- House price prediction

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41	880	129.0	322	126	8.3252	452600	NEAR BAY
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3014	358500	NEAR BAY
2	-122.24	37.85	52	1467	190.0	496	177	7.2574	352100	NEAR BAY
3	-122.25	37.85	52	1274	235.0	558	219	5.6431	341300	NEAR BAY
4	-122.25	37.85	52	1627	280.0	565	259	3.8462	342200	NEAR BAY

# Practical Step 1

- 1. Data preprocessing

```
df = pd.read_csv('housing.csv')

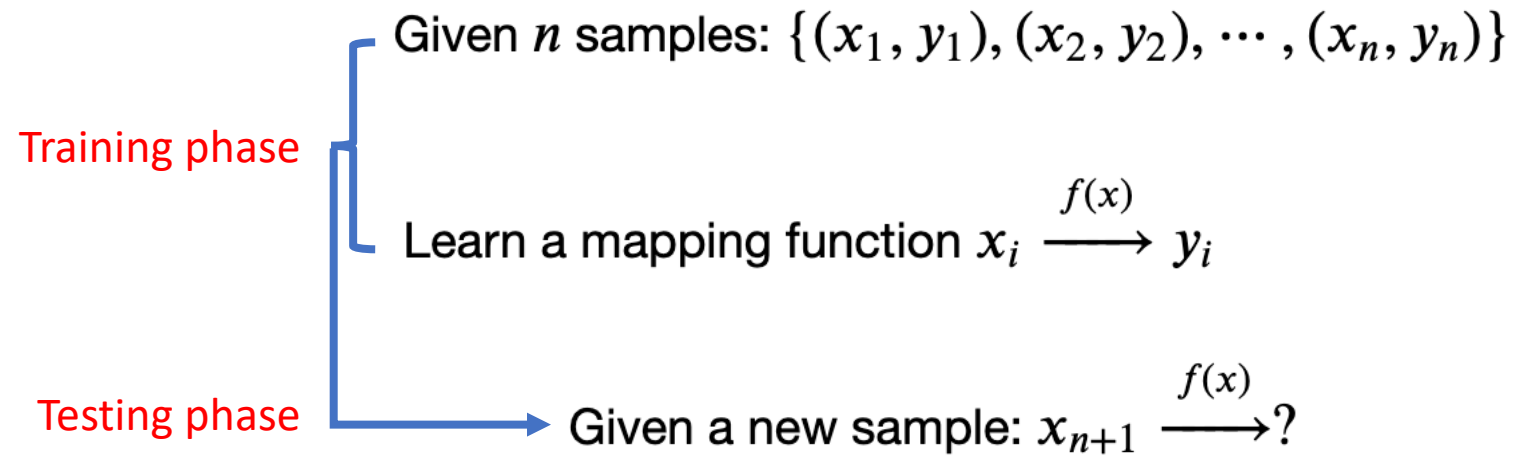
# 0. fill in missing values
mean_val = df['total_bedrooms'].mean()
df['total_bedrooms'] = df['total_bedrooms'].fillna(mean_val)
print(df.isnull().sum())

# 1. convert categorical features to numerical values
labelencoder = LabelEncoder()
df['ocean_proximity'] = labelencoder.fit_transform(df['ocean_proximity'])
print(df.info())
```



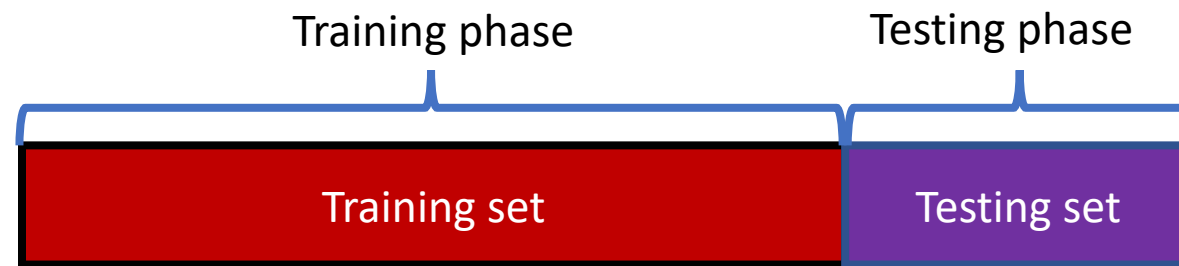
# Practical Step 2

- 2. Split dataset



# Practical Step 2

- 2. Split dataset
  - Training set
    - Used for training the model, **during the training phase**
  - Testing set
    - Used for evaluating the performance of the model, **after obtaining the model**



# Practical Step 2

- *# 2. split samples*

```
house_fea = df.drop('median_house_value', axis=1).values
```

```
house_price = df['median_house_value'].values
```

```
X_train, X_test, y_train, y_test = train_test_split(house_fea,  
                                                    house_price,  
                                                    test_size=0.2,  
                                                    random_state=42)
```

```
print(X_train.shape)      (16512, 9)
```

```
print(X_test.shape)       (4128, 9)
```

- *# normalize features*

```
normalizer = StandardScaler()
```

```
X_train = normalizer.fit_transform(X_train)
```

```
X_test = normalizer.transform(X_test)
```

# Practical Step 3

- 3. Train the model
  - Find the optimal model parameter

```
#3. train the model  
lr = LinearRegression()  
  
lr.fit(X_train,y_train)  
  
print("bias is "+str(lr.intercept_))  
print("coefficients is "+str(lr.coef_))
```

```
bias is 207194.69373786208  
coefficients is [-85854.94724101 -90946.06271148  14924.30655143 -17693.23405277  
 48767.60670995 -43884.16852449  17601.31495096  77144.10164179  
 -451.52015229]
```

# Practical Step 3

- 3. Train the model

```
y_train_pred = lr.predict(X_train)

mae = mean_absolute_error(y_train_pred, y_train)
mse = mean_squared_error(y_train_pred, y_train)
rmse = np.sqrt(mse)

print('prediction for training set:')
print('MAE is: {}'.format(mae))
print('MSE is: {}'.format(mse))
print('RMSE is: {}'.format(rmse))
```

```
prediction for training set:
MAE is: 50626.9285430206
MSE is: 4810958229.787787
RMSE is: 69361.0714290645
```

# Practical Step 4

- 4. Evaluate the model

```
#4. evaluate the model
y_test_pred = lr.predict(X_test)


mae = mean_absolute_error(y_test_pred, y_test)
mse = mean_squared_error(y_test_pred, y_test)
rmse = np.sqrt(mse)

print('prediction for testing set:')
print('MAE is: {}'.format(mae))
print('MSE is: {}'.format(mse))
print('RMSE is: {}'.format(rmse))
```

```
prediction for testing set:
MAE is: 51846.87784903816
MSE is: 5055025116.165613
RMSE is: 71098.69982050033
```

# Practical Step 4

- 4. Evaluate the model




```
labels = ['House1', 'House2', 'House3', 'House4', 'House5']
x = np.arange(len(labels)) # the label locations
width = 0.35 # the width of the bars

fig, ax = plt.subplots()
rects1 = ax.bar(x - width/2, y_test[0:5], width, label='ground truth')
rects2 = ax.bar(x + width/2, y_test_pred[0:5], width, label='prediction')


ax.set_ylabel('Price')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()
```

# How to fit non-linear data?

- Y has a non-linear response to X



Non-linear model



Linear model



# How to fit non-linear data?

- Apply non-linear transformation to features

One-dimensional input:  $\mathbf{x} = (x) \xrightarrow{\phi(\mathbf{x})} \{1, x, x^2, x^3, \dots, x^k\}$

$$f(x) = w_0 + w_1 x + w_2 x^2 + w_3 x^3 + \dots + w_k x^k$$

Two-dimensional input:  $\mathbf{x} = (x_1, x_2) \xrightarrow{\phi(\mathbf{x})} \{1, x_1, x_1^2, x_2, x_2^2, x_1 x_2\}$

$$f(x) = w_0 + w_1 x_1 + w_2 x_1^2 + w_3 x_2 + w_4 x_2^2 + w_5 x_1 x_2$$

# How to fit non-linear data?

- Linear model

x	y
1	5.14
2	10.67
3	16.17
4	20.12

$$f(x_1) = w_0 + w_1 * 1$$

$$f(x_2) = w_0 + w_1 * 2$$

$$f(x_3) = w_0 + w_1 * 3$$

$$f(x_4) = w_0 + w_1 * 4$$

- Non-linear model

x	x^2	x^3	x^4	y
1	1^2	1^3	1^4	5.14
2	2^2	2^3	2^4	10.67
3	3^2	3^3	3^4	16.17
4	4^2	4^3	4^4	20.12

$$f(x_1) = w_0 + w_1 * 1 + w_2 * 1^2 + w_3 * 1^3 + w_4 * 1^4$$

$$f(x_2) = w_0 + w_1 * 2 + w_2 * 2^2 + w_3 * 2^3 + w_4 * 2^4$$

$$f(x_3) = w_0 + w_1 * 3 + w_2 * 3^2 + w_3 * 3^3 + w_4 * 3^4$$

$$f(x_4) = w_0 + w_1 * 4 + w_2 * 4^2 + w_3 * 4^3 + w_4 * 4^4$$

# How to fit non-linear data?

- Linear model

x_1	x_2	y
1	-1	11
2	-2	12
3	-3	13
4	-4	14





- Non-linear model Two-dimensional input:  $\mathbf{x} = (x_1, x_2) \xrightarrow{\phi(\mathbf{x})} \{1, x_1, x_1^2, x_2, x_2^2, x_1 x_2\}$

x_1	x_2				y
1	-1				11
2	-2				12
3	-3				13
4	-4				14

# How to fit non-linear data?


- Non-linear transformation  $\phi(x) = \{1, x, x^2, \dots, x^k\}$

$$f(x) = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_kx^k$$



# Overfitting

- Errors on training data are small
- But errors on **new points** are likely to be large



# How to avoid overfitting?

- Add a regularization term
  - Make some  $w_i$  very small or approach to zero

$$f(x) = w_0 + w_1x + w_2x^2 + w_3x^3 + \cdots + w_kx^k$$

$$\min_w \boxed{\frac{1}{n} \|y - Xw\|_2^2} + \boxed{\lambda \|w\|_2^2}$$

Fit the data

Control the model complexity

$w$  is the model parameter

$\lambda$  is the hyperparameter

# How to avoid overfitting?

- Optimization

$$\min_w \frac{1}{n} \|y - Xw\|_2^2 + \lambda \|w\|_2^2$$




$$\frac{\partial \mathcal{L}(w)}{\partial w} = -\frac{2}{n} X^T (y - Xw) + 2\lambda w = 0$$



$$(X^T X + \lambda n I) w = X^T y$$



$$w = (X^T X + \lambda n I)^{-1} X^T y$$



$$\mathcal{L}(w) = \frac{1}{n} \|y - Xw\|_2^2$$



$$\frac{\partial \mathcal{L}(w)}{\partial w} = -\frac{2}{n} X^T (y - Xw) = 0$$



$$X^T X w = X^T y$$



$$w = (X^T X)^{-1} X^T y$$

# Ridge Regression vs Lasso

- Ridge regression
  - $\ell_2$ -norm regularization

$$\min_w \frac{1}{n} \|y - Xw\|_2^2 + \lambda \|w\|_2^2$$

- Lasso
  - $\ell_1$ -norm regularization

$$\min_w \frac{1}{n} \|y - Xw\|_2^2 + \lambda \|w\|_1$$



# Ridge Regression vs Lasso

- Ridge Regression

```
#3. train the model  
lr = Ridge(alpha=0.1)  
  
lr.fit(X_train,y_train)
```

- Lasso

```
#3. train the model  
lr = Lasso(alpha=0.1)  
  
lr.fit(X_train,y_train)
```

# Ridge Regression vs Lasso

```
x = np.random.rand(10, 1)
noise = np.random.rand(10, 1)*0.2
y = 3*x + noise

# linear regression
lr = LinearRegression()
lr.fit(x,y)
print("coefficients is "+str(lr.coef_))

# ridge regression
lr = Ridge(alpha=0.1)
lr.fit(x,y)
print("coefficients is "+str(lr.coef_))

# lasso
lr = Lasso(alpha=0.1)
lr.fit(x,y)
print("coefficients is "+str(lr.coef_))
```

```
coefficients is [[2.95186796]]
coefficients is [[2.64284605]]
coefficients is [1.78259097]
```

# Online Resources

- [https://cims.nyu.edu/~cfgranda/pages/OBDA\\_fall17/notes/linear\\_models.pdf](https://cims.nyu.edu/~cfgranda/pages/OBDA_fall17/notes/linear_models.pdf)
- <https://stat.ethz.ch/education/semesters/ss2016/regression/Regression.pdf>
- [https://www.cs.toronto.edu/~rgrosse/courses/csc311\\_f20/readings/notes on linear regression.pdf](https://www.cs.toronto.edu/~rgrosse/courses/csc311_f20/readings/notes_on_linear_regression.pdf)