

Logistic Regression

Hongchang Gao

Spring 2024

Linear Regression vs Logistic Regression

Given n samples: $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

- Linear regression

- Y is continuous
- Model

$$f(x_i) = w_0 + w_1 x_{i,1} + w_2 x_{i,2} + \dots + w_d x_{i,d}$$

- Loss function

$$\min_w \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$$

- Logistic regression

- Y is discrete
- Model

$$f(\mathbf{x}_i) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x}_i)}$$

- Loss function

$$\min_{\mathbf{w}} - \sum_{i=1}^n \{y_i \log p_i + (1 - y_i) \log(1 - p_i)\}$$

Evaluation for Binary Classification

- When the data is **balanced**, use the classification accuracy
 - Balanced data: #positive samples is almost the same with #negative samples

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

Ground-truth	Prediction	Correct?
1	0	N
1	1	Y
0	0	Y
0	0	Y
1	1	Y
1	1	Y
0	1	N

Accuracy=5/7

Evaluation for Binary Classification

- When the data is **imbalanced**, does accuracy still work?
 - Imbalanced data: #positive samples is very different from #negative samples

Ground-truth	Prediction	Correct?
1	0	N
0	0	Y
0	0	Y
0	0	Y
0	0	Y
0	0	Y
0	0	Y

Accuracy=6/7

1. The classifier gives the same prediction for all samples,
But it still has a large accuracy!

2. Accuracy is not a good metric for imbalanced data!

Evaluation for Binary Classification

- Example
 - COVID-19 diagnosis: a good classifier should find all positive cases.
 - Fraud detection: a good classifier should find all fraud transactions.

Ground-truth	Prediction	Correct?
1	0	N
0	0	Y
0	0	Y
0	0	Y
0	0	Y
0	0	Y
0	0	Y

1. This classifier fails to find positive cases or fraud transactions
2. But it still has a large accuracy

Evaluation for Binary Classification

- True Positive (TP)
 - Prediction is positive, ground-truth is positive
- False Positive (FP)
 - Prediction is positive, ground-truth is negative
- False Negative (FN)
 - Prediction is negative, ground-truth is positive
- True Negative (TN)
 - Prediction is negative, ground-truth is negative

	Positive (ground-truth)	Negative (ground-truth)
Positive (prediction)	True Positive (TP)	False Positive (FP)
Negative (prediction)	False Negative (FN)	True Negative (TN)

Evaluation for Binary Classification

- Example: COVID-19 diagnosis

Ground-truth	Prediction
1	1
0	0
1	0
0	1
0	0
0	1
0	0
0	0

	Positive (ground-truth)	Negative (ground-truth)
Positive (prediction)	1 (TP)	2 (FP)
Negative (prediction)	1 (FN)	4 (TN)

Evaluation for Binary Classification

- When the data is **imbalanced**, use recall or precision
 - **Recall**: the proportion of actual positives is classified correctly

$$\text{Recall} = \frac{TP}{TP+FN}$$

- **Precision**: the proportion of positive predictions is actually correct

$$\text{Precision} = \frac{TP}{TP+FP}$$

	Positive (ground-truth)	Negative (ground-truth)
Positive (prediction)	True Positive (TP)	False Positive (FP)
Negative (prediction)	False Negative (FN)	True Negative (TN)

Evaluation for Binary Classification

- Example

Ground-truth	Prediction
1	1
0	0
1	0
0	1
0	0
0	1
0	0
0	0

	Positive (ground-truth)	Negative (ground-truth)
Positive (prediction)	1 (TP)	2 (FP)
Negative (prediction)	1 (FN)	4 (TN)

$$\text{Recall} = 1/2$$

$$\text{Precision} = 1/3$$

$$\text{Accuracy} = 5/8$$

$$\text{Recall} = \frac{TP}{TP+FN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

Evaluation for Binary Classification

- Example

	Positive (ground-truth)	Negative (ground-truth)
Positive (predict)	1	1
Negative (predict)	8	90

Accuracy = $(90+1)/100=0.91$ But for the positive class, only one sample is classified correctly

Recall = $1/(1+8)=0.11$ It only correctly recognizes 11% of all positive samples


Precision = $1/(1+1)=0.5$ When the classifier gives a positive prediction, it is correct 50% of the time

Evaluation for Binary Classification

- When to care about **recall**?
 - It is important to find all positive samples
 - e.g. COVID-19 diagnosis, fraud detection
 - A larger recall means more ground-truth positive samples were found by the classifier
- When to care about **precision**?
 - Users are sensitive to the prediction error
 - e.g. google search
 - A larger precision means the positive prediction has a large probability to be correct.

Evaluation for Binary Classification

- Relationship between recall and precision



- F1 score: harmonic mean of precision and recall
 - conveys the balance between the precision and the recall

$$F_1 = \frac{2 \times \text{Recall} \times \text{Precision}}{\text{Recall} + \text{Precision}}$$

Summary

- Build model
 - Linear classifier to separate positive and negative samples

$$\min_{\mathbf{w}} \sum_{i=1}^n \{\log(1 + \exp(\mathbf{w}^T \mathbf{x}_i)) - y_i \mathbf{w}^T \mathbf{x}_i\} + \lambda \|\mathbf{w}\|_2^2$$

- Optimize the model

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \eta \sum_{i=1}^n \mathbf{x}_i (y_i - \frac{\mathbf{w}_k^T \mathbf{x}_i}{1 + \exp(\mathbf{w}_k^T \mathbf{x}_i)}) - 2\eta \lambda \mathbf{w}_k$$

- Evaluate the model
 - Balanced: Accuracy
 - Imbalanced: Recall, Precision, F1 score

Practical Steps

- 1. Preprocess data

```
# preprocess the dataset
X, y = datasets.load_iris(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.1,
                                                    random_state=0)

# normalize the dataset
normalizer = StandardScaler()
X_train = normalizer.fit_transform(X_train)
X_test = normalizer.transform(X_test)
```

Practical Steps

- 2. Train the model

```
# train logistic regression
clf = LogisticRegression(penalty='l2', C=1.0)
clf.fit(X_train, y_train)

y_train_pred = clf.predict(X_train)
acc = accuracy_score(y_train, y_train_pred)
print("Training accuracy: {:.4f}".format(acc))
```

Training accuracy: 0.9259

Practical Steps

- 3. Evaluate the model

```
# evaluate logistic regression  
y_test_pred = clf.predict(X_test)  
acc = accuracy_score(y_test, y_test_pred)  
print("Testing accuracy: {:.4f}".format(acc))
```


Testing accuracy: 0.8000

Model Selection

- Model selection
 - Different hyperparameters lead to different models/performance

$$\min_{\mathbf{w}} \sum_{i=1}^n \{\log(1 + \exp(\mathbf{w}^T \mathbf{x}_i)) - y_i \mathbf{w}^T \mathbf{x}_i\} + \lambda \|\mathbf{w}\|_2^2$$

- How to find a model with the best performance?



```
regularization_coefficient = [0.1, 0.5, 1.0, 5.0, 10.0]

for reg in regularization_coefficient:
    clf = LogisticRegression(penalty='l2', C=reg)

    clf.fit(X_train, y_train)

    y_test_pred = clf.predict(X_test)


    acc = accuracy_score(y_test, y_test_pred)

    print("reg_coeff: {}, acc: {:.3f}".format(1.0/reg, acc))
```

```
reg_coeff: 10.0, acc: 0.533
reg_coeff: 2.0, acc: 0.533
reg_coeff: 1.0, acc: 0.533
reg_coeff: 0.2, acc: 0.800
reg_coeff: 0.1, acc: 0.800
```

Model Selection: Threefold Split

- Threefold split:
 - Training set
 - Used for training the model, during the training phase
 - Validation set
 - Used for hyperparameter selection, during the training phase
 - Testing set
 - Used for evaluating the model, after obtaining the model



Do NOT use the testing set to select the hyperparameter

Model Selection: Threefold Split

- How to select the model (hyperparameters)?
 - For each hyperparameter
 - Train the model using the training set
 - Evaluate the model using the validation set
 - Use the performance on the validation set to select the best hyperparameters
- After obtaining the best hyperparameter:
 - Retrain the model with both the training and validation sets
 - Report the performance on the testing set

Model Selection: Threefold Split

```
X, y = datasets.load_iris(return_X_y=True)
print(X.shape)
```

```
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y,
                                                            test_size=0.1,
                                                            random_state=0)
print(X_train_val.shape, X_test.shape)
```

```
X_train, X_valid, y_train, y_valid = train_test_split(X_train_val,
                                                       y_train_val,
                                                       test_size=0.1,
                                                       random_state=0)
print(X_train.shape, X_valid.shape)
```

(150, 4)	
(135, 4)	(15, 4)
(121, 4)	(14, 4)

Model Selection: Threefold Split

```
regularization_coefficient = [0.1, 0.5, 1.0, 5.0, 10.0]

best_acc = 0.0
best_reg = 0.0

for reg in regularization_coefficient:
    clf = LogisticRegression(penalty='l2', C=reg)

    clf.fit(X_train, y_train) Using the training set

    y_valid_pred = clf.predict(X_valid) Using the validation set
    acc = accuracy_score(y_valid, y_valid_pred)

    if acc > best_acc:
        best_acc = acc
        best_reg = reg

    print("reg_coeff: {}, acc: {:.3f}".format(1.0/reg, acc))
```

Model selection based on the validation set

```
reg_coeff: 10.0, acc: 0.571
reg_coeff: 2.0, acc: 0.643
reg_coeff: 1.0, acc: 0.643
reg_coeff: 0.2, acc: 0.714
reg_coeff: 0.1, acc: 0.714
acc: 0.800
```

```
clf = LogisticRegression(penalty='l2', C=best_reg)
clf.fit(X_train_val, y_train_val)
y_test_pred = clf.predict(X_test)
acc = accuracy_score(y_test, y_test_pred)

print("acc: {:.3f}".format(acc))
```

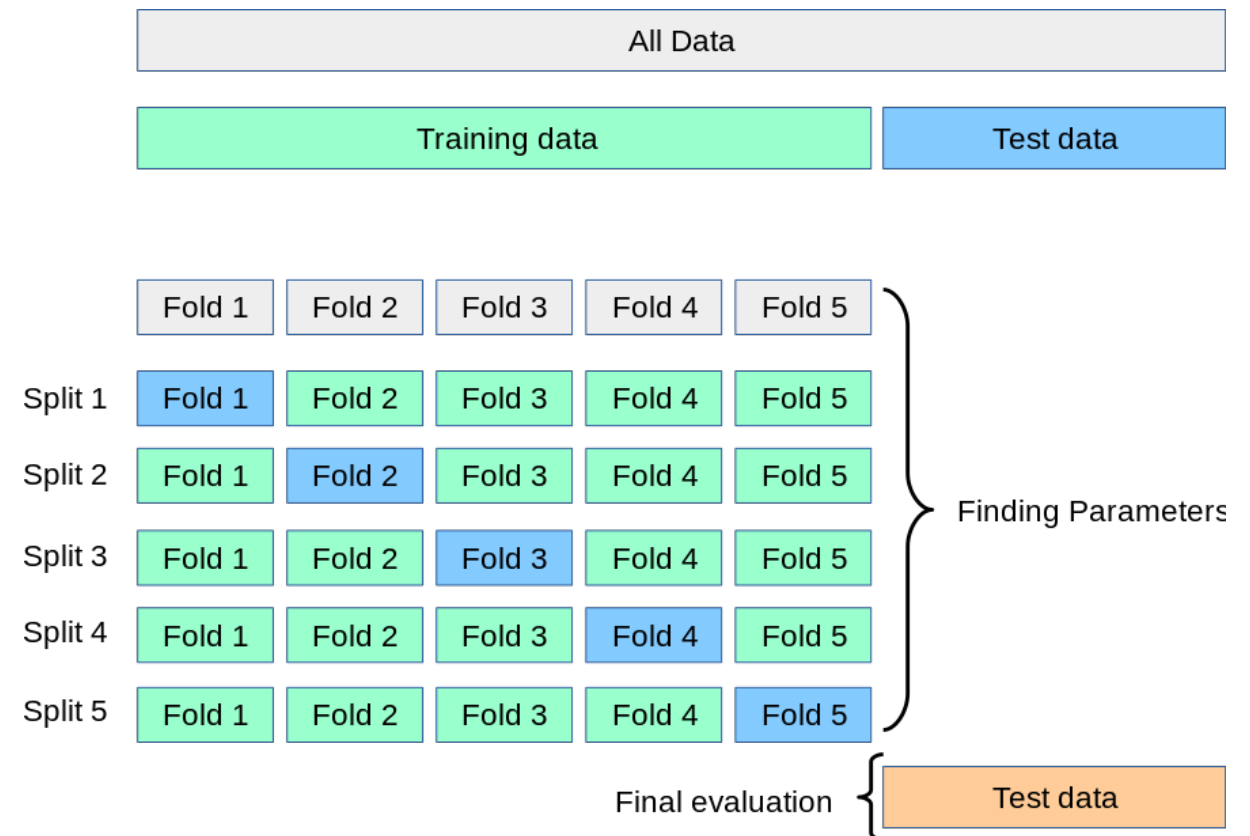
Retrain the model with the best hyperparameter,
Evaluate the model on the testing set

Model Selection: Threefold Split

- Pros:
 - Fast and simple
- Cons:
 - Large variance
 - Unreliable estimation of future performance
 - Bad use of data
 - Waste data

Model Selection: Cross-validation

- Training data:
 - Randomly partition it into **K** folds
 - K-1 folds for **training set**
 - 1 fold for **validation set**
- How to select the model?
 - For each hyperparameter
 - Train the model for K times
 - Evaluate the model for K times
 - Use the mean of K evaluation to select model



```
for reg in regularization_coefficient:
```

```
    sum_acc = 0.0
```

```
    for fold in range(5):
```

```
        index_of_folds_temp = index_of_folds.copy()
```

```
        valid_index = index_of_folds_temp[fold,:].reshape(-1)
```

```
        train_index = np.delete(index_of_folds_temp, fold, 0).reshape(-1)
```

```
        X_train = X_train_val[train_index]
```

```
        y_train = y_train_val[train_index]
```

```
        X_valid = X_train_val[valid_index]
```

```
        y_valid = y_train_val[valid_index]
```

```
        clf = LogisticRegression(penalty='l2', C=reg)
```

```
        clf.fit(X_train, y_train)
```

```
        y_valid_pred = clf.predict(X_valid)
```

```
        acc = accuracy_score(y_valid, y_valid_pred)
```

```
        sum_acc += acc
```

```
    cur_acc = sum_acc / 5.0
```

```
    print("reg_coeff: {}, acc: {:.3f}".format(1.0/reg, cur_acc))
```

```
    if cur_acc > best_acc:
```

```
        best_acc = cur_acc
```

```
        best_reg = reg
```

n

reg_coeff: 10.0, acc: 0.770

reg_coeff: 2.0, acc: 0.911

reg_coeff: 1.0, acc: 0.941

reg_coeff: 0.2, acc: 0.948

reg_coeff: 0.1, acc: 0.948

Model Selection: Cross-validation

```
clf = LogisticRegression(penalty='l2', C=best_reg)
clf.fit(X_train_val, y_train_val)

y_test_pred = clf.predict(X_test)
acc = accuracy_score(y_test, y_test_pred)

print("acc: {:.3f}".format(acc))
```

acc: 0.867

Model Selection: Cross-validation

- Pros:
 - More data
 - More stable
- Cons:
 - Slower